
Table of Contents

1. Introduction
2. Data Loading and Merging
3. Data Preprocessing I
 - 3.1 Convert Timestamp
 - 3.2 Creating New Features
 - 3.3 Checking Duplicates
 - 3.4 Remove Unnecessary Columns
 - 3.5 Handling Missing Values
 - 3.6 Outlier Detection
 - 3.7 Data Quality Issues
 - 3.8 Summary of Findings
4. Exploratory Data Analysis (EDA)
 - 4.1 Basic Statistics
 - 4.2 Univariate Analysis
 - 4.3 Bivariate Analysis
 - 4.4 Multivariate Analysis
5. Model Training and Evaluation
6. Results Summary and Visualization

1. Introduction

Our analysis aims to understand and predict parking availability across 57 distinct roads over a four-month period. The dataset includes information on 13,439,520 parking slots and captures various influencing factors such as weather conditions, transportation options, and nearby facilities like restaurants and schools. By analyzing these variables, we seek to identify patterns and key determinants of parking availability. This analysis will provide actionable insights for optimizing parking management through dynamic pricing, improving infrastructure, and enhancing the overall user experience. Understanding these dynamics will help us develop effective strategies to ensure better parking availability for users.

2. Data Loading and Merging

- Loading datasets
- Displaying Sample Data
- Merging three datasets

```
# Importing required libraries
import pandas as pd
import numpy as np
import seaborn as sns
```

```

import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
from sklearn.linear_model import LinearRegression

# Load datasets
ground_df = pd.read_csv('groundtruth.csv')
weather_df = pd.read_csv('weather_features.csv')
road_df = pd.read_csv('road_features.csv')

ground_df.head()

    Unnamed: 0  road_segment_id          timestamp  max_capacity
occupied \
0           0            20026  2019-03-01 06:00:00             7
3
1           1            20026  2019-03-01 06:09:00             7
4
2           2            20026  2019-03-01 06:11:00             7
4
3           3            20026  2019-03-01 06:19:00             7
3
4           4            20026  2019-03-01 06:20:00             7
3

    available
0           4
1           3
2           3
3           4
4           4

weather_df.head()

    Unnamed: 0  road_segment_id          timestamp  tempC
windspeedKmph \
0           0            20026  2019-03-01 06:00:00        24
18
1           1            20026  2019-03-01 06:09:00        24
18
2           2            20026  2019-03-01 06:11:00        24
18
3           3            20026  2019-03-01 06:19:00        24
18
4           4            20026  2019-03-01 06:20:00        24
18

```

```

    precipMM
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0

road_df.head()

   Unnamed: 0  road_segment_id  commercial  residential
transportation \
0             0            20026       58.0        32
0
1           21960            20057       92.0        14
0
2           43920            20069       35.0         5
1
3           65880            20190       33.0         7
1
4           87840            20198       73.0         4
1

schools  eventsites  restaurant  shopping  office  supermarket \
0        4           2          100       100       65        63
1        0           2          100       100      100        66
2        0           1          100       100       43        21
3        0           1          100       100       65        26
4        0           2          100       100       40        10

  num_off_street_parking  off_street_capa
0                      4                 462
1                      4                 335
2                      4                 442
3                      7                 719
4                      4                 277

# Merge groundtruth with weather_features on road_segment_id and timestamp
merged_df = pd.merge(ground_df, weather_df, on=['road_segment_id', 'timestamp'])

# Merge the result with road_features on road_segment_id
df = pd.merge(merged_df, road_df, on='road_segment_id')

df.head()

   Unnamed: 0_x  road_segment_id  timestamp  max_capacity
occupied \
0             0            20026  2019-03-01 06:00:00            7
3

```

1	1	20026	2019-03-01	06:09:00	7
4	2	20026	2019-03-01	06:11:00	7
4	3	20026	2019-03-01	06:19:00	7
3	4	20026	2019-03-01	06:20:00	7
3					
available Unnamed: 0_y tempC windspeedKmph precipMM ...					
residential \					
0	4	0	24	18	0.0 ...
32					
1	3	1	24	18	0.0 ...
32					
2	3	2	24	18	0.0 ...
32					
3	4	3	24	18	0.0 ...
32					
4	4	4	24	18	0.0 ...
32					
transportation schools eventsites restaurant shopping					
office \					
0	0	4	2	100	100 65
1	0	4	2	100	100 65
2	0	4	2	100	100 65
3	0	4	2	100	100 65
4	0	4	2	100	100 65
supermarket num_off_street_parking off_street_capa					
0	63		4	462	
1	63		4	462	
2	63		4	462	
3	63		4	462	
4	63		4	462	
[5 rows x 22 columns]					

3. Data Preprocessing I

- Convert timestamp
- Extract few columns
- Checking duplicates

- Remove unnecessary columns
- Checking missing values
- Checking outliers
- Checking data quality issues

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1251720 entries, 0 to 1251719
Data columns (total 22 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Unnamed: 0_x      1251720 non-null   int64  
 1   road_segment_id  1251720 non-null   int64  
 2   timestamp         1251720 non-null   object  
 3   max_capacity     1251720 non-null   int64  
 4   occupied          1251720 non-null   int64  
 5   available         1251720 non-null   int64  
 6   Unnamed: 0_y      1251720 non-null   int64  
 7   tempC             1251720 non-null   int64  
 8   windspeedKmph    1251720 non-null   int64  
 9   precipMM          1251720 non-null   float64 
 10  Unnamed: 0        1251720 non-null   int64  
 11  commercial        1229760 non-null   float64 
 12  residential       1251720 non-null   int64  
 13  transportation    1251720 non-null   int64  
 14  schools            1251720 non-null   int64  
 15  eventsites        1251720 non-null   int64  
 16  restaurant         1251720 non-null   int64  
 17  shopping           1251720 non-null   int64  
 18  office              1251720 non-null   int64  
 19  supermarket        1251720 non-null   int64  
 20  num_off_street_parking 1251720 non-null   int64  
 21  off_street_capa   1251720 non-null   int64  
dtypes: float64(2), int64(19), object(1)
memory usage: 210.1+ MB
```

3.1 Convert Timestamp

```
# Convert timestamp
df['timestamp'] = pd.to_datetime(df['timestamp'])
```

3.2 Creating New Features

```
# Extract the date part of the timestamp
df['date'] = df['timestamp'].dt.date

# Extract the time part of the timestamp
df['time'] = df['timestamp'].dt.time
```

3.3 Checing Duplicates

```
# Check for duplicates
duplicates = df.duplicated(subset=['road_segment_id', 'timestamp'])
print(f"Number of duplicate rows: {duplicates.sum()}")

Number of duplicate rows: 0
```

3.4 Remove unnecessary columns

```
# Remove the unnecessary columns
df = df.drop(['Unnamed: 0_x', 'Unnamed: 0_y', 'Unnamed: 0'], axis =
'columns')

df.shape


(1251720, 21)
```

3.5 Handling Missing Values

```
# Checking the null values
df.isnull().sum()

road_segment_id          0
timestamp                0
max_capacity             0
occupied                 0
available                0
tempC                     0
windspeedKmph             0
precipMM                 0
commercial               21960
residential              0
transportation            0
schools                  0
eventsites                0
restaurant                0
shopping                  0
office                    0
supermarket                0
num_off_street_parking      0
off_street_capa            0
date                      0
time                      0
dtype: int64


# Visualize missing values
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title('Missing Values Heatmap')
plt.show()
```



3.6 Outlier Detection


```
# Boxplot for numerical features to detect outliers
numerical_features = ['max_capacity', 'occupied', 'available',
'tempC', 'windspeedKmph', 'precipMM',
'commercial', 'residential', 'transportation',
'schools', 'eventsites',
'restaurant', 'shopping', 'office',
'supermarket', 'num_off_street_parking', 'off_street_capa']

plt.figure(figsize=(15, 10))
df[numerical_features].boxplot()
plt.xticks(rotation=90)
plt.title('Boxplot for Numerical Features')
plt.show()
```




```
# Detailed boxplot for each feature
for feature in numerical_features:
    plt.figure(figsize=(10, 5))
    sns.boxplot(x=df[feature])
    plt.title(f'Boxplot for {feature}')
    plt.show()
```


Boxplot for max_capacity




Boxplot for occupied




Boxplot for available




Boxplot for tempC




Boxplot for windspeedKmph




Boxplot for precipMM




Boxplot for commercial




Boxplot for residential




Boxplot for transportation




Boxplot for schools




Boxplot for eventsites




Boxplot for restaurant




Boxplot for shopping




Boxplot for office




Boxplot for supermarket



Boxplot for num_off_street_parking



Boxplot for off_street_capa



```
def count_outliers(df):

    outliers = {}
    total_points = df.shape[0]

    for column in df.select_dtypes(include=['float64',
'int64']).columns:
        # Calculate Q1 and Q3
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)

        # Calculate IQR
        IQR = Q3 - Q1

        # Define outlier boundaries
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        # Count outliers
        num_outliers = df[(df[column] < lower_bound) | (df[column] >
upper_bound)].shape[0]
        percentage_outliers = (num_outliers / total_points) * 100

        outliers[column] = [num_outliers, percentage_outliers]

    # Convert dictionary to df
    outliers_df = pd.DataFrame.from_dict(outliers, orient='index',
```

```
columns=['Number of Outliers', 'Percentage of Outliers'])
```

```
    return outliers_df
```

```
outliers_df = count_outliers(df)
outliers_df
```

	Number of Outliers	Percentage of Outliers
road_segment_id	109800	8.771930
max_capacity	109800	8.771930
occupied	97081	7.755808
available	160035	12.785208
tempC	11629	0.929042
windspeedKmph	2052	0.163934
precipMM	284544	22.732240
commercial	0	0.000000
residential	0	0.000000
transportation	219600	17.543860
schools	21960	1.754386
eventsites	0	0.000000
restaurant	0	0.000000
shopping	0	0.000000
office	131760	10.526316
supermarket	87840	7.017544
num_off_street_parking	0	0.000000
off_street_capa	65880	5.263158

3.7 Data Quality issues

```
# Identifying negative values
```

	road_segment_id	timestamp	max_capacity	
occupied	7425	20026 2019-04-11 09:46:00	7	8
	8178	20026 2019-04-15 12:32:00	7	8

12103	20026	2019-05-07	09:38:00	7	8
13917	20026	2019-05-17	10:45:00	7	9
16318	20026	2019-05-30	15:54:00	7	8
...
1228924	22942	2019-06-26	11:21:00	2	3
1228966	22942	2019-06-26	14:52:00	2	3
1229101	22942	2019-06-27	11:08:00	2	3
1229165	22942	2019-06-27	16:27:00	2	3
1229354	22942	2019-06-28	17:13:00	2	3
available tempC windspeedKmph precipMM commercial					
residential \					
7425	-1	14	4	0.0	58.0
32					
8178	-1	23	13	0.0	58.0
32					
12103	-1	13	15	0.0	58.0
32					
13917	-2	14	7	0.0	58.0
32					
16318	-1	11	18	0.0	58.0
32					
...
...					
1228924	-1	12	8	0.0	20.0
10					
1228966	-1	13	8	0.0	20.0
10					
1229101	-1	13	17	0.0	20.0
10					
1229165	-1	13	15	0.0	20.0
10					
1229354	-1	13	20	0.0	20.0
10					
... schools eventsites restaurant shopping office					
supermarket \					
7425	...	4	2	100	100
63					
8178	...	4	2	100	100

63						
12103	...	4	2	100	100	65
63						
13917	...	4	2	100	100	65
63						
16318	...	4	2	100	100	65
63						
...
1228924	...	0	1	100	100	15
22						
1228966	...	0	1	100	100	15
22						
1229101	...	0	1	100	100	15
22						
1229165	...	0	1	100	100	15
22						
1229354	...	0	1	100	100	15
22						
		num_off_street_parking	off_street_capa		date	time
7425		4		462	2019-04-11	09:46:00
8178		4		462	2019-04-15	12:32:00
12103		4		462	2019-05-07	09:38:00
13917		4		462	2019-05-17	10:45:00
16318		4		462	2019-05-30	15:54:00
...	
1228924		5		1133	2019-06-26	11:21:00
1228966		5		1133	2019-06-26	14:52:00
1229101		5		1133	2019-06-27	11:08:00
1229165		5		1133	2019-06-27	16:27:00
1229354		5		1133	2019-06-28	17:13:00

[1353 rows x 21 columns]

```
def count_negative_values(df, columns):
    negative_values = {}
    total_points = df.shape[0]
```

```

    for column in columns:
        # Count negative values
        num_negative_values = (df[column] < 0).sum()
        percentage_negative_values = (num_negative_values /
total_points) * 100

            negative_values[column] = [num_negative_values,
percentage_negative_values]

    # Convert dictionary to df
    negative_values_df = pd.DataFrame.from_dict(negative_values,
orient='index', columns=['Number of Negative Values', 'Percentage of
Negative Values'])

    return negative_values_df

```

negative_values_df = count_negative_values(df, count_based_columns)
negative_values_df

	Number of Negative Values \
max_capacity	0
occupied	0
available	1353
windspeedKmph	0
precipMM	0
commercial	0
residential	0
transportation	0
schools	0
eventsites	0
restaurant	0
shopping	0
office	0
supermarket	0
num_off_street_parking	0
off_street_capa	0

	Percentage of Negative Values
max_capacity	0.000000
occupied	0.000000
available	0.108091
windspeedKmph	0.000000
precipMM	0.000000
commercial	0.000000
residential	0.000000
transportation	0.000000
schools	0.000000
eventsites	0.000000
restaurant	0.000000
shopping	0.000000

office	0.000000
supermarket	0.000000
num_off_street_parking	0.000000
off_street_capa	0.000000

3.8 Summary of Findings

```
data_issues = [
    {
        "Problem/Logical Issue": "Missing values",
        "Affected Variable/s": "commercial",
        "Number of Observations": 21960,
        "Potential Issues": "Missing values can lead to inaccurate analysis and model predictions.",
        "Handling Technique/s": "Imputation techniques such as filling missing values with mean, median, or using interpolation."
    },
    {
        "Problem/Logical Issue": "Outliers",
        "Affected Variable/s": "Available, precipMM, transportation, off_street_capa",
        "Number of Observations": "Various (see above)",
        "Potential Issues": "Outliers can distort the model training process and lead to poor predictions.",
        "Handling Technique/s": "Identifying and removing outliers using statistical methods like the cap value method."
    },
    {
        "Problem/Logical Issue": "Duplicates",
        "Affected Variable/s": "None",
        "Number of Observations": 0,
        "Potential Issues": "Multiple entries for the same road_segment_id and timestamp can skew results.",
        "Handling Technique/s": "Identifying and removing duplicate records to ensure each observation is unique."
    },
    {
        "Problem/Logical Issue": "Data quality issue 1",
        "Affected Variable/s": "occupied, available, max_capacity",
        "Number of Observations": "available: 1353",
        "Potential Issues": "sum of occupied and available should always equal max capacity. As a result of that we found 1,353 negative values in available.",
        "Handling Technique/s": "Replaced by 0"
    },
    {
        "Problem/Logical Issue": "Data quality issue 2",
        "Affected Variable/s": "tempC",
        "Number of Observations": "1",
    }
]
```

```

        "Potential Issues": "mistake",
        "Handling Technique/s": "Replaced by 0/ mean"
    }
]

df_issues = pd.DataFrame(data_issues)
df_issues

  Problem/Logical Issue                               Affected
Variable/s \
0           Missing values
commercial
1           Outliers Available, precipMM, transportation,
off_street...
2           Duplicates
None
3 Data quality issue 1                         occupied, available,
max_capacity
4 Data quality issue 2
tempC

  Number of Observations                           Potential
Issues \
0                      21960 Missing values can lead to inaccurate
analysis...
1 Various (see above)   Outliers can distort the model training
process...
2                         0 Multiple entries for the same
road_segment_id ...
3 available: 1353 sum of occupied and available should always
eq...
4                         1
mistake

  Handling Technique/s
0 Imputation techniques such as filling missing ...
1 Identifying and removing outliers using statis...
2 Identifying and removing duplicate records to ...
3                                     Replaced by 0
4                                     Replaced by 0/ mean

```

4. Exploratory Data Analysis (EDA)

- Basic statistics
- Univariate Analysis
- Bivariate Analysis
- Multivariate Analysis

4.1 Basic Statistics

```
pd.set_option('display.float_format', '{:.2f}'.format)

# Basic statistics
df.describe()

    road_segment_id           timestamp  max_capacity \
count      1251720.00          1251720      1251720.00
mean       22097.02  2019-05-01 01:29:30.076503040      10.74
min        20026.00          2019-03-01 06:00:00      1.00
25%        21943.00          2019-03-31 13:31:15      3.00
50%        22277.00          2019-05-01 01:31:00      7.00
75%        22622.00          2019-05-31 13:26:30     12.00
max        23277.00          2019-06-30 20:59:00     66.00
std         738.07                  NaN      12.66

    occupied  available   tempC  windspeedKmph  precipMM
commercial \
count  1251720.00  1251720.00  1251720.00      1251720.00  1251720.00
1229760.00
mean       2.41       8.32     15.63       12.75      0.10
23.14
min        0.00      -5.00    -999.00       1.00      0.00
0.00
25%        0.00       2.00     12.00       8.00      0.00
2.00
50%        1.00       4.00     15.00      12.00      0.00
10.00
75%        3.00       8.00     19.00      17.00      0.00
48.00
max        34.00      66.00     32.00      32.00      4.80
92.00
std        3.90      12.15      5.10       6.43      0.37
26.91

    residential  transportation  schools  eventsites  restaurant
\
count      1251720.00          1251720.00  1251720.00  1251720.00  1251720.00
mean       54.70          0.79       4.00       1.00      69.53
min        2.00          0.00       0.00       0.00      18.00
25%        7.00          0.00       0.00       0.00      32.00
50%       18.00          0.00       1.00       1.00     100.00
75%      119.00          1.00       9.00       2.00     100.00
max       175.00          3.00      28.00       4.00     100.00
```

```

std          60.37        1.15       5.33      1.08      33.62

      shopping    office supermarket num_off_street_parking \
count 1251720.00 1251720.00   1251720.00           1251720.00
mean   54.82     15.70     13.30            1.23
min    5.00      1.00      0.00            0.00
25%    11.00     2.00      2.00            0.00
50%    35.00     5.00      5.00            0.00
75%   100.00    20.00     20.00            3.00
max   100.00    100.00    66.00            7.00
std   42.04     21.07    17.15            1.82

      off_street_capa
count      1251720.00
mean      161.51
min       0.00
25%       0.00
50%       0.00
75%     277.00
max    1133.00
std    265.22

df['max_capacity'].sum()
13439520

df['road_segment_id'].nunique()
57

df.shape
(1251720, 21)

```


4.2 Univariate Analysis

```

# Histograms for each numeric column
def plot_histograms(df):
    df.hist(bins=10, figsize=(15, 10), color='skyblue',
edgecolor='black')
    plt.suptitle('Univariate Analysis - Histograms of Numeric
Features')
    plt.tight_layout()
    plt.show()

plot_histograms(df)


```



4.3 Bivariate Analysis

```
# Scatter plot matrix
def plot_scatter_matrix(df, features):
    sns.pairplot(df[features], plot_kws={'alpha':0.5, 's':30})
    plt.suptitle('Bivariate Analysis - Scatter Plot Matrix', y=1.02)
    plt.show()

plot_scatter_matrix(df, ['max_capacity', 'occupied', 'available',
                           'tempC', 'windspeedKmph'])
```




```
# Traffic Flow and Capacity Utilization
# Plot daily occupancy trends to identify peak and low times
plt.figure(figsize=(14, 6))
plt.plot(df['date'], df['occupied'], label='Occupied Spaces',
color='blue')
plt.plot(df['date'], df['available'], label='Available Spaces',
color='orange')
plt.xlabel('Date')
plt.ylabel('Number of Spaces')
plt.title('Traffic Flow - Daily Occupancy Trends')
plt.legend()
```

```

plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

C:\Users\lahir\AppData\Local\Temp\ipykernel_17740\307938723.py:11:
UserWarning: Creating legend with loc="best" can be slow with large
amounts of data.
    plt.tight_layout()
C:\Users\lahir\anaconda3\Lib\site-packages\IPython\core\
pylabtools.py:170: UserWarning: Creating legend with loc="best" can be
slow with large amounts of data.
    fig.canvas.print_figure(bytes_io, **kw)

```



Cant interpret the above graph and therefore grouping the data


```

# Daily Average Occupancy Trends
# Grouping data by day to calculate daily averages of occupied and
available spaces
df['date'] = df['timestamp'].dt.date
daily_avg = df.groupby('date').agg({'occupied': 'mean', 'available':
'mean'})

# Plot daily trends with the aggregated data
plt.figure(figsize=(14, 6))
plt.plot(daily_avg.index, daily_avg['occupied'], label='Occupied
Spaces', color='blue')
plt.plot(daily_avg.index, daily_avg['available'], label='Available
Spaces', color='orange')
plt.xlabel('Date')
plt.ylabel('Average Number of Spaces')
plt.title('Daily Average Occupancy Trends')
plt.legend()
plt.xticks(rotation=45)

```

```
plt.tight_layout()
plt.show()
```



```
# Temporally removing the outlier
df2= df[df['tempC'] != -999]


# Environmental Impact on Usage
# Scatter plot of occupancy vs temperature, wind speed, and precipitation
fig, axs = plt.subplots(1, 3, figsize=(18, 5))

sns.scatterplot(x=df2['tempC'], y=df2['occupied'], ax=axs[0],
color='coral')
axs[0].set_title('Occupancy vs Temperature (°C)')
axs[0].set_xlabel('Temperature (°C)')
axs[0].set_ylabel('Occupied Spaces')

sns.scatterplot(x=df['windspeedKmph'], y=df['occupied'], ax=axs[1],
color='skyblue')
axs[1].set_title('Occupancy vs Wind Speed (Kmph)')
axs[1].set_xlabel('Wind Speed (Kmph)')
axs[1].set_ylabel('Occupied Spaces')


sns.scatterplot(x=df['precipMM'], y=df['occupied'], ax=axs[2],
color='green')
axs[2].set_title('Occupancy vs Precipitation (mm)')
axs[2].set_xlabel('Precipitation (mm)')
axs[2].set_ylabel('Occupied Spaces')

plt.suptitle('Environmental Impact on Occupancy')
plt.tight_layout()
plt.show()
```



```
# Off-street Parking Capacity vs Occupancy
plt.figure(figsize=(14, 6))
plt.plot(df['date'], df['num_off_street_parking'], label='Actual Off-Street Usage', color='purple')
plt.plot(df['date'], df['off_street_capa'], label='Off-Street Capacity', color='gray', linestyle='--')
plt.xlabel('Date')
plt.ylabel('Number of Off-Street Parking Spaces')
plt.title('Off-Street Parking Capacity vs Actual Usage')
plt.legend()
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

C:\Users\lahir\AppData\Local\Temp\ipykernel_17740\2958170429.py:10:
UserWarning: Creating legend with loc="best" can be slow with large
amounts of data.
plt.tight_layout()
C:\Users\lahir\anaconda3\Lib\site-packages\IPython\core\pylabtools.py:170: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
fig.canvas.print_figure(bytes_io, **kw)




```

# Off-Street Parking Capacity vs Actual Usage
# Aggregate data by date (for example, by summing up the
# 'num_off_street_parking' and 'off_street_capa')
daily_data = df.groupby('date').agg({
    'num_off_street_parking': 'sum',
    'off_street_capa': 'sum'
}).reset_index()

# Plot the aggregated data
plt.figure(figsize=(14, 6))
plt.plot(daily_data['date'], daily_data['num_off_street_parking'],
label='Actual Off-Street Usage', color='purple')
plt.plot(daily_data['date'], daily_data['off_street_capa'],
label='Off-Street Capacity', color='gray', linestyle='--')

plt.xlabel('Date')
plt.ylabel('Number of Off-Street Parking Spaces')
plt.title('Off-Street Parking Capacity vs Actual Usage')
plt.legend(loc='upper left')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



```

# Check the unique date counts
date_counts = df['date'].value_counts()
print("Unique date counts:\n", date_counts)

Unique date counts:
date
2019-03-01    10260
2019-05-31    10260
2019-05-29    10260
2019-05-28    10260

```

```

2019-05-27    10260
...
2019-04-06    10260
2019-04-05    10260
2019-04-04    10260
2019-04-03    10260
2019-06-30    10260
Name: count, Length: 122, dtype: int64

# Define the time range for events
# Assume events occur between 18:00 and 20:59
start_time = pd.to_datetime('18:00:00').time()
end_time = pd.to_datetime('20:59:00').time()

# Filter the DataFrame for times between 18:00 and 20:59
df_events = df[(df['time'] >= start_time) & (df['time'] <= end_time)]
df_events

   occupied  road_segment_id      timestamp  max_capacity
144          \  20026  2019-03-01 18:03:00                7         2
145          \  20026  2019-03-01 18:05:00                7         2
146          \  20026  2019-03-01 18:12:00                7         0
147          \  20026  2019-03-01 18:18:00                7         0
148          \  20026  2019-03-01 18:22:00                7         1
...
1251715        \  23277  2019-06-30 20:36:00                9         5
1251716        \  23277  2019-06-30 20:44:00                9         5
1251717        \  23277  2019-06-30 20:46:00                9         5
1251718        \  23277  2019-06-30 20:52:00                9         5
1251719        \  23277  2019-06-30 20:59:00                9         5

   residential  available  tempC  windspeedKmph  precipMM  commercial
144            \       5     29                 9      0.00      58.00
32
145            \       5     29                 9      0.00      58.00
32
146            \       7     29                 9      0.00      58.00
32

```

147	7	29	9	0.00	58.00	
32						
148	6	29	9	0.00	58.00	
32						
...	
...						
1251715	4	9	18	0.00	11.00	
16						
1251716	4	9	18	0.00	11.00	
16						
1251717	4	9	18	0.00	11.00	
16						
1251718	4	9	18	0.00	11.00	
16						
1251719	4	9	18	0.00	11.00	
16						
supermarket	...	schools	eventsites	restaurant	shopping	office
144	...	4	2	100	100	65
63						
145	...	4	2	100	100	65
63						
146	...	4	2	100	100	65
63						
147	...	4	2	100	100	65
63						
148	...	4	2	100	100	65
63						
...
...						
1251715	...	1	1	100	100	2
7						
1251716	...	1	1	100	100	2
7						
1251717	...	1	1	100	100	2
7						
1251718	...	1	1	100	100	2
7						
1251719	...	1	1	100	100	2
7						
		num_off_street_parking	off_street_capa		date	time
144		4		462	2019-03-01	18:03:00
145		4		462	2019-03-01	18:05:00
146		4		462	2019-03-01	18:12:00

147	4	462	2019-03-01	18:18:00
148	4	462	2019-03-01	18:22:00
...
1251715	0	0	2019-06-30	20:36:00
1251716	0	0	2019-06-30	20:44:00
1251717	0	0	2019-06-30	20:46:00
1251718	0	0	2019-06-30	20:52:00
1251719	0	0	2019-06-30	20:59:00

[250344 rows x 21 columns]

```
# Event Impacts and Temporal Patterns Analysis

# Create a column to simulate events; assume events occur during
# specific hours across days
df['is_event_time'] = df['time'].apply(lambda t: 1 if t.hour in [18,
19, 20] else 0) # Events in the evening

# Average Occupancy During Event vs Non-Event Times
# Calculate average occupancy for event and non-event times
event_vs_non_event = df.groupby('is_event_time')['occupied'].mean()
event_vs_non_event.index = ['Non-Event Times', 'Event Times']


# Plot event vs non-event time occupancy
plt.figure(figsize=(8, 6))
sns.barplot(x=event_vs_non_event.index, y=event_vs_non_event.values,
palette='pastel')
plt.xlabel('Event Time')
plt.ylabel('Average Occupied Spaces')
plt.title('Average Occupancy During Event vs Non-Event Times')
plt.show()
```

C:\Users\lahir\AppData\Local\Temp\ipykernel_17740\1450093152.py:8:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=event_vs_non_event.index, y=event_vs_non_event.values,
palette='pastel')
```

Average Occupancy During Event vs Non-Event Times



```
# Average Occupancy by Day of the Week
# Convert 'timestamp' to datetime
df['timestamp'] = pd.to_datetime(df['timestamp'], errors='coerce')

# Temporal Patterns: Weekly and Daily Trends
# Extract day of the week and hour of the day for temporal pattern
# analysis
df['day_of_week'] = df['timestamp'].dt.day_name()
df['hour_of_day'] = df['timestamp'].dt.hour

# Average occupancy by day of the week
avg_occupancy_by_day = df.groupby('day_of_week')
['occupied'].mean().reindex(
    ["Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
     "Saturday", "Sunday"])
)


# Plot
plt.figure(figsize=(10, 6))
sns.barplot(x=avg_occupancy_by_day.index,
```

```
y=avg_occupancy_by_day.values, palette='viridis')
plt.xlabel('Day of the Week')
plt.ylabel('Average Occupied Spaces')
plt.title('Average Occupancy by Day of the Week')
plt.show()
```

```
C:\Users\lahir\AppData\Local\Temp\ipykernel_17740\3574020592.py:17:
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.barplot(x=avg_occupancy_by_day.index,
y=avg_occupancy_by_day.values, palette='viridis')
```




```
# Average occupancy by hour of the day
avg_occupancy_by_hour = df.groupby('hour_of_day')['occupied'].mean()

# Plot average occupancy by hour of the day
plt.figure(figsize=(12, 6))
sns.lineplot(x=avg_occupancy_by_hour.index,
y=avg_occupancy_by_hour.values, marker='o', color='teal')
plt.xlabel('Hour of the Day')
plt.ylabel('Average Occupied Spaces')
```

```

plt.title('Average Occupancy by Hour of the Day')
plt.grid()
plt.tight_layout()
plt.show()

```




```

# Calculate average occupancy and available spaces by hour of the day
avg_occupancy_by_hour = df.groupby('hour_of_day')['occupied'].mean()
avg_available_by_hour = df.groupby('hour_of_day')['available'].mean()

# Plot average occupancy and available spaces by hour of the day
plt.figure(figsize=(12, 6))
sns.lineplot(x=avg_occupancy_by_hour.index,
y=avg_occupancy_by_hour.values, marker='o', color='teal',
label='Average Occupied Spaces')
sns.lineplot(x=avg_available_by_hour.index,
y=avg_available_by_hour.values, marker='o', color='orange',
label='Average Available Spaces')


# Add labels and title
plt.xlabel('Hour of the Day')
plt.ylabel('Average Spaces')
plt.title('Average Occupancy and Available Spaces by Hour of the Day')
plt.grid()
plt.legend()
plt.tight_layout()
plt.show()

```



```
# Monthly Average Occupancy Trends
df['month'] = df['timestamp'].dt.month
monthly_avg = df.groupby('month').agg({'occupied': 'mean',
                                         'available': 'mean'})
month_labels = {3: 'Mar', 4: 'Apr', 5: 'May', 6: 'Jun'}

plt.figure(figsize=(10, 6))
plt.plot(monthly_avg.index, monthly_avg['occupied'], label='Occupied Spaces', color='blue')
# plt.plot(monthly_avg.index, monthly_avg['available'],
#          label='Available Spaces', color='orange')
plt.xlabel('Month')
plt.ylabel('Average Number of Spaces')
plt.title('Monthly Average Occupancy Trends')
plt.legend()
plt.xticks(monthly_avg.index, [month_labels[m] for m in monthly_avg.index])
plt.tight_layout()
plt.show()
```




```

# Daily Occupancy Ratio Trends
# Calculate occupancy to availability ratio
df['occupancy_ratio'] = df['occupied'] / (df['occupied'] +
df['available'])

# Group by day and calculate daily occupancy ratio
daily_ratio = df.groupby('date')['occupancy_ratio'].mean()


# Plot the daily occupancy ratio trend
plt.figure(figsize=(14, 6))
plt.plot(daily_ratio.index, daily_ratio.values, label='Occupancy
Ratio', color='purple')
plt.xlabel('Date')
plt.ylabel('Occupancy Ratio')
plt.title('Daily Occupancy Ratio Trends')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



```
# Average Occupancy under Different Weather Conditions
# Divide data into high and low temperature/wind groups
high_temp = df[df['tempC'] > df['tempC'].median()]
low_temp = df[df['tempC'] <= df['tempC'].median()]
high_wind = df[df['precipMM'] > df['precipMM'].median()]
low_wind = df[df['precipMM'] <= df['precipMM'].median()]

# Plot comparison of occupancy under different conditions
plt.figure(figsize=(10, 6))
plt.bar(['High Temp', 'Low Temp', 'High precipMM', 'Low precipMM'],
        [high_temp['occupied'].mean(), low_temp['occupied'].mean(),
         high_wind['occupied'].mean(), low_wind['occupied'].mean()],
        color=['salmon', 'skyblue', 'orange', 'purple'])
plt.ylabel('Average Occupied Spaces')
plt.title('Average Occupancy under Different Weather Conditions')
plt.show()
```




```
df.columns

Index(['road_segment_id', 'timestamp', 'max_capacity', 'occupied',
'available',
       'tempC', 'windspeedKmph', 'precipMM', 'commercial',
'residential',
       'transportation', 'schools', 'eventsites', 'restaurant',
'shopping',
       'office', 'supermarket', 'num_off_street_parking',
'off_street_capa',
       'date', 'time', 'is_event_time', 'day_of_week', 'hour_of_day',
'month',
       'occupancy_ratio'],
      dtype='object')

## 4.4 Multivariate Analysis


# Heatmap for correlation
def plot_correlation_heatmap(df):
    corr = df.corr()
    plt.figure(figsize=(10, 8))
    sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
    plt.title('Multivariate Analysis - Correlation Heatmap')
    plt.show()

plot_correlation_heatmap(df.iloc[:, 2:-7]) # Considering only the
relevant features
```



```
# Analyzing the impact of nearby facilities and off-street parking on occupancy


# Correlation heatmap
plt.figure(figsize=(10, 8))
facilities_cols = ['commercial', 'residential', 'transportation',
'schools', 'eventsites',
'restaurant', 'shopping', 'office', 'supermarket']
correlation_matrix = df[['occupied']] + facilities_cols].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1,
vmax=1)
plt.title('Correlation Between Occupancy and Nearby Facilities')
plt.show()
```



```
# Heatmap of Average Hourly Occupancy by Day of the Week
# Create a pivot table to show average occupancy per hour for each day of the week
df['day_of_week'] = df['timestamp'].dt.dayofweek
df['hour'] = df['timestamp'].dt.hour
hourly_weekday_pivot = df.pivot_table(values='occupied', index='day_of_week', columns='hour', aggfunc='mean')

# Plot heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(hourly_weekday_pivot, cmap='YlGnBu', annot=True, fmt=".1f")
plt.xlabel('Hour of the Day')
plt.ylabel('Day of the Week')
plt.title('Heatmap of Average Hourly Occupancy by Day of the Week')
plt.yticks(ticks=range(7), labels=['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])
```


```
'Sat', 'Sun'], rotation=0)
plt.show()
```



```
# Temporal Analysis
# Line plot for occupied vs available over time
def plot_temporal_trends(df):
    plt.figure(figsize=(12, 6))
    plt.plot(df['date'], df['occupied'], label='Occupied',
color='blue')
    plt.plot(df['date'], df['available'], label='Available',
color='orange')
    plt.xlabel('Date')
    plt.ylabel('Count')
    plt.title('Temporal Analysis - Occupied vs Available Over Date')
    plt.legend()
    plt.xticks(rotation=45)
    plt.show()

plot_temporal_trends(df)
```

```
C:\Users\lahir\anaconda3\Lib\site-packages\IPython\core\
pylabtools.py:170: UserWarning: Creating legend with loc="best" can be
slow with large amounts of data.
fig.canvas.print_figure(bytes_io, **kw)
```



Cluster Analysis

```

# Prepare data for clustering by averaging occupancy at each hour of
# the day
hourly_pattern = df.groupby('hour_of_day').agg({'occupied': 'mean',
'available': 'mean'})

# Standardize data for clustering
scaler = StandardScaler()
hourly_pattern_scaled = scaler.fit_transform(hourly_pattern)

# Apply k-means clustering
kmeans = KMeans(n_clusters=3,
random_state=0).fit(hourly_pattern_scaled)
hourly_pattern['cluster'] = kmeans.labels_

# Plot hourly patterns by cluster
plt.figure(figsize=(10, 6))
for cluster in range(3):
    subset = hourly_pattern[hourly_pattern['cluster'] == cluster]
    plt.plot(subset.index, subset['occupied'], label=f'Cluster {cluster}', marker='o')

plt.xlabel('Hour of the Day')
plt.ylabel('Average Occupied Spaces')
plt.title('Clustered Occupancy Patterns by Hour')
plt.legend()


```

```

plt.xticks(range(0, 24))
plt.tight_layout()
plt.show()

C:\Users\lahir\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1446: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
warnings.warn(

```



5. Model Training and Evaluation

I ran several experiments with different data preprocessing techniques, as shown in the results graph below. I conducted the experiments without feature engineering, by selecting specific features, by creating new features using lag features and rolling aggregates for max capacity and occupancy, and finally, by performing feature engineering and treating outliers.

```

# Placeholder for results tracking
results = []

# Define a function to train and evaluate a model
def run_experiment(df, features, target, attempt_name, top_n=10):
    # Splitting the data
    X = df[features]

```

```

y = df[target]
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=False)

# Normalization using StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Model training
model = RandomForestRegressor()
model.fit(X_train, y_train)

# Predictions
predictions = model.predict(X_test)

# Calculate evaluation metrics
mae = mean_absolute_error(y_test, predictions)
rmse = np.sqrt(mean_squared_error(y_test, predictions))
r2 = r2_score(y_test, predictions)

# Print metrics
print(f"Attempt: {attempt_name}")
print(f"Mean Absolute Error: {mae}")
print(f"Root Mean Squared Error: {rmse}")
print(f"R-squared: {r2}")

# Append metrics to results list
results.append({
    'Attempt': attempt_name,
    'MAE': mae,
    'RMSE': rmse,
    'R2': r2
})

# Get feature importances
importances = model.feature_importances_
importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

# Select the top N most important features
top_features = importance_df.head(top_n)

# Plot top important features
plt.figure(figsize=(10, 6))
plt.barh(top_features['Feature'], top_features['Importance'],
color='skyblue')
plt.xlabel('Feature Importance')

```


```

plt.ylabel('Features')
plt.title(f'Top {top_n} Important Features for {attempt_name}')
plt.gca().invert_yaxis()
plt.show()

# Attempt 1 - Selecting all the features and without doing feature engineering
attempt1 = ['max_capacity', 'occupied', 'tempC',
            'windspeedKmph', 'precipMM', 'commercial', 'residential',
            'transportation', 'schools', 'eventsites', 'restaurant',
            'shopping',
            'office', 'supermarket', 'num_off_street_parking',
            'off_street_capa']
run_experiment(df, attempt1, 'available', "Attempt 1")

Attempt: Attempt 1
Mean Absolute Error: 1.0882860783561816
Root Mean Squared Error: 2.089467644245773
R-squared: 0.9151142918809467

```



```

# Attempt 2 - Feature Engineering, Preprocessing and selecting specific features

# Extract useful temporal features from the date and time columns
df['timestamp'] = pd.to_datetime(df['timestamp'])
df['day_of_week'] = df['timestamp'].dt.dayofweek
df['is_weekend'] = df['day_of_week'].apply(lambda x: 1 if x >= 5 else 0)

```

```

df['hour'] = pd.to_datetime(df['timestamp']).dt.hour

# Drop unnecessary columns
#df.drop(columns=['date', 'time'], inplace=True)

df['is_weekend'] = df['day_of_week'].isin([5, 6]).astype(int)
df['hour'] = pd.to_datetime(df['timestamp']).dt.hour

# Replace missing values in 'commercial' column with the median
df['commercial'].fillna(df['commercial'].median(), inplace=True)

# Replace missing or invalid values in 'available' with 0
df['available'].fillna(0, inplace=True)
df['available'] = df['available'].apply(lambda x: 0 if x < 0 else x)

# Replace -999 in 'tempC' column with the average of the column
mean_tempC = df.loc[df['tempC'] != -999, 'tempC'].mean() # Calculate mean without -999 values
df['tempC'].replace(-999, mean_tempC, inplace=True)

attempt2 = ['tempC', 'windspeedKmph', 'precipMM', 'commercial',
'residential',
                     'transportation', 'schools', 'eventsites',
'restaurant', 'shopping',
                     'office', 'supermarket',
'num_off_street_parking', 'off_street_capa',
                     'day_of_week', 'is_weekend', 'hour']
run_experiment(df, attempt2, 'available', "Attempt 2")

```

C:\Users\lahir\AppData\Local\Temp\ipykernel_17740\1126651581.py:16:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

df['commercial'].fillna(df['commercial'].median(), inplace=True)
C:\Users\lahir\AppData\Local\Temp\ipykernel_17740\1126651581.py:19:  

FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

```


For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['available'].fillna(0, inplace=True)
C:\Users\lahir\AppData\Local\Temp\ipykernel_17740\1126651581.py:24:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['tempC'].replace(-999, mean_tempC, inplace=True)

Attempt: Attempt 2
Mean Absolute Error: 7.331413768012332
Root Mean Squared Error: 8.661392299071043
R-squared: -0.4586998063368364
```




```
# Attempt 3 - Feature Engineering, Preprocessing and selecting all features
attempt3 = ['max_capacity', 'occupied','tempC', 'windspeedKmph',
'precipMM', 'commercial',
'residential', 'transportation', 'schools', 'eventsites',
'restaurant', 'shopping',
'office', 'supermarket',
'num_off_street_parking', 'off_street_capa',
'day_of_week', 'is_weekend', 'hour']
run_experiment(df, attempt3, 'available', "Attempt 3")
```

Attempt: Attempt 3

Mean Absolute Error: 1.1459261655961397

Root Mean Squared Error: 2.2659800618889285

R-squared: 0.9001604362218294



```
# Attempt 4 - More feature engineering (lag and rolling features :
occupied)
df['occupied_lag_1'] = df['occupied'].shift(1)
df['occupied_lag_3'] = df['occupied'].shift(3)
df['occupied_lag_6'] = df['occupied'].shift(6)
df['occupied_rolling_mean_3'] =
df['occupied'].rolling(window=3).mean()
df['occupied_rolling_mean_6'] =
df['occupied'].rolling(window=6).mean()
df['occupied_rolling_mean_12'] =
df['occupied'].rolling(window=12).mean()
df.fillna(0, inplace=True)
```

```

attempt4 = ['max_capacity', 'occupied', 'occupied_lag_1',
'occupied_lag_3', 'occupied_rolling_mean_3',
'occupied_rolling_mean_6', 'tempC',
'windspeedKmph', 'precipMM', 'commercial',
'residential', 'transportation', 'schools',
'eventsites', 'restaurant',
'shopping', 'office', 'supermarket',
'num_off_street_parking', 'off_street_capa',
'day_of_week', 'is_weekend', 'hour']
run_experiment(df, attempt4, 'available', "Attempt 4")


```

Attempt: Attempt 4

Mean Absolute Error: 1.1653860288243378

Root Mean Squared Error: 2.299110173079787

R-squared: 0.8972196546641568



```

# Attempt 5 - More feature engineering (lag and rolling features :
occupied, max_capacity)
df['max_capacity_lag_1'] = df['max_capacity'].shift(1)
df['max_capacity_lag_3'] = df['max_capacity'].shift(3)
df['max_capacity_lag_6'] = df['max_capacity'].shift(6)
df['max_capacity_rolling_mean_3'] =
df['max_capacity'].rolling(window=3).mean()
df['max_capacity_rolling_mean_6'] =
df['max_capacity'].rolling(window=6).mean()
df['max_capacity_rolling_mean_12'] =
df['max_capacity'].rolling(window=12).mean()

```


```

df.fillna(0, inplace=True)

attempt5 = ['max_capacity', 'occupied', 'occupied_lag_1',
'occupied_lag_3', 'occupied_rolling_mean_3',
'occupied_rolling_mean_6', 'max_capacity_lag_1',
'max_capacity_lag_3', 'max_capacity_rolling_mean_3',
'max_capacity_rolling_mean_6', 'tempC',
'windspeedKmph', 'precipMM', 'commercial',
'residential', 'transportation', 'schools',
'eventsites', 'restaurant',
'shopping', 'office', 'supermarket',
'num_off_street_parking', 'off_street_capa',
'day_of_week', 'is_weekend', 'hour']
run_experiment(df, attempt5, 'available', "Attempt 5")

```

Attempt: Attempt 5
Mean Absolute Error: 1.014532722973189
Root Mean Squared Error: 2.0527633303156416
R-squared: 0.9180652244334556



```

# Attempt 6 - More feature engineering (lag and rolling features:  

# occupied, max_capacity) and treating outliers  

# Capping outliers at the 99th percentile for precipMM and  

# transportation
precip_cap = df['precipMM'].quantile(0.99)
trans_cap = df['transportation'].quantile(0.99)

df['precipMM'] = np.where(df['precipMM'] > precip_cap, precip_cap,
df['precipMM'])

```

```

df['transportation'] = np.where(df['transportation'] > trans_cap,
trans_cap, df['transportation'])

attempt6 = ['max_capacity', 'occupied', 'occupied_lag_1',
'occupied_lag_3', 'occupied_rolling_mean_3',
'occupied_rolling_mean_6', 'max_capacity_lag_1',
'max_capacity_lag_3', 'max_capacity_rolling_mean_3',
'max_capacity_rolling_mean_6', 'tempC',
'windspeedKmph', 'precipMM', 'commercial',
'residential', 'transportation', 'schools',
'eventsites', 'restaurant',
'shopping', 'office', 'supermarket',
'num_off_street_parking', 'off_street_capa',
'day_of_week', 'is_weekend', 'hour']
run_experiment(df, attempt6, 'available', "Attempt 6")


```

Attempt: Attempt 6

Mean Absolute Error: 1.0984482152558084

Root Mean Squared Error: 2.2488428796709456

R-squared: 0.9016648615763506



6. Results Summary and Visualization

```

# Convert results to DataFrame
results_df = pd.DataFrame(results)

# Mapping old attempt names to new names
attempt_names = {

```

```

'Attempt 1': 'W/0 FE or Preprocessing',
'Attempt 2': 'FE & Specific Features',
'Attempt 3': 'FE & All Features',
'Attempt 4': 'More FE',
'Attempt 5': 'More FE 2',
'Attempt 6': 'More FE & Outliers',
}

# Rename attempts
results_df['Attempt'] = results_df['Attempt'].map(attempt_names)
results_df

      Attempt    MAE    RMSE     R2
0  W/0 FE or Preprocessing  1.09   2.09  0.92
1  FE & Specific Features  7.33   8.66 -0.46
2  FE & All Features       1.15   2.27  0.90
3  More FE                  1.17   2.30  0.90
4  More FE 2                1.01   2.05  0.92
5  More FE & Outliers      1.10   2.25  0.90

# Plotting the results
fig, ax1 = plt.subplots(figsize=(12, 6))

# Plot MAE, RMSE, and R-squared on the same plot
ax1.plot(results_df['Attempt'], results_df['MAE'], marker='o',
label='MAE', color='blue')
ax1.plot(results_df['Attempt'], results_df['RMSE'], marker='o',
label='RMSE', color='orange')
ax2 = ax1.twinx()
ax2.plot(results_df['Attempt'], results_df['R2'], marker='o',
label='R-squared', color='green')

# Labeling the plots
ax1.set_xlabel("Attempts")
ax1.set_ylabel("MAE / RMSE")
ax2.set_ylabel("R-squared")
ax1.legend(loc="upper left")
ax2.legend(loc="upper right")
plt.title("Model Evaluation Metrics Across Different Attempts")
plt.savefig("evaluation_metrics.png")
plt.show()

```

